

Name: Bernardo, Christian Emmanuel	Date Performed: 02/04/2024
Course/Section: CPE 232-CPE31S1	Date Submitted: 02/04/2024
Instructor: Dr. Jonathan Taylar	Semester and SY: 2nd, 2023-2024
Activity 10: Install, Configure, and Manage Log Monitoring tools	
1. Objectives	
Create and design a workflow that installs, configure and manage enterprise log monitoring tools using Ansible as an Infrastructure as Code (IaC) tool.	
2. Discussion	
<p>Log monitoring software scans and monitors log files generated by servers, applications, and networks. By detecting and alerting users to patterns in these log files, log monitoring software helps solve performance and security issues. System administrators use log monitoring software to detect common important events indicated by log files.</p> <p>Log monitoring software helps maintain IT infrastructure performance and pinpoints issues to prevent downtime and mitigate risks. These tools will often integrate with IT alerting software, log analysis software, and other IT issue resolution products to more aptly flesh out the IT infrastructure maintenance ecosystem.</p> <p>To qualify for inclusion in the Log Monitoring category, a product must:</p> <ul style="list-style-type: none"> • Monitor the log files generated by servers, applications, or networks • Alert users when important events are detected • Provide reporting capabilities for log files <p>Elastic Stack</p> <p>ELK suite stands for Elasticsearch, Kibana, Beats, and Logstash (also known as the ELK Stack). Source: https://www.elastic.co/elastic-stack</p> <p>The Elastic Stack is a group of open source products from Elastic designed to help users take data from any type of source and in any format, and search, analyze and visualize that data in real time. The product group was formerly known as the ELK Stack for the core products in the group -- Elasticsearch, Logstash and Kibana -- but has been rebranded as the Elastic Stack. A fourth product, Beats, was subsequently added to the stack. The Elastic Stack can be deployed on premises or made available as software as a service (SaaS). Elasticsearch supports Amazon Web Services (AWS), Google Cloud Platform and Microsoft Azure.</p> <p>GrayLog</p>	

Graylog is a powerful platform that allows for easy log management of both structured and unstructured data along with debugging applications.

It is based on Elasticsearch, MongoDB, and Scala. Graylog has a main server, which receives data from its clients installed on different servers, and a web interface, which visualizes the data and allows to work with logs aggregated by the main server.

We use Graylog primarily as the stash for the logs of the web applications we build. However, it is also effective when working with raw strings (i.e. syslog): the tool parses it into the structured data we need. It also allows advanced custom search in the logs using structured queries. In other words, when integrated properly with a web app, Graylog helps engineers to analyze the system behavior on almost per code line basis.

Source: <https://www.graylog.org/products/open-source>

3. Tasks

1. Create a playbook that:
 - a. Install and configure Elastic Stack in separate hosts (Elastic Search, Kibana, Logstash)
2. Apply the concept of creating roles.
3. Describe how you did step 1. (Provide screenshots and explanations in your report. Make your report detailed such that it will look like a manual.)
4. Show an output of the installed Elastic Stack for both Ubuntu and CentOS.
5. Make sure to create a new repository in GitHub for this activity.

4. Output (screenshots and explanations)


step one: Setting up the repository for this activity

Create a new repository

A repository contains all project files, including the revision history. Already have a project repository elsewhere? [Import a repository.](#)

Required fields are marked with an asterisk (*).

Owner *

 Xerxes000

Repository name *

CPE232_Bernardo_ACT-10.1

✔ CPE232_Bernardo_ACT-10.1 is available.

Great repository names are short and memorable. Need inspiration? How about [upgraded-journey](#) ?

Description (optional)

☒  **Public**

Anyone on the internet can see this repository. You choose who can commit.

☐  **Private**

You choose who can see and commit to this repository.

Initialize this repository with:

☒ **Add a README file**

This is where you can write a long description for your project. [Learn more about READMEs.](#)

```
christian@workstation: ~  
christian@workstation:~$ git clone git@github.com:Xerxes000/CPE232_Bernardo_ACT-10.1.git  
Cloning into 'CPE232_Bernardo_ACT-10.1'...  
remote: Enumerating objects: 3, done.  
remote: Counting objects: 100% (3/3), done.  
remote: Total 3 (delta 0), reused 0 (delta 0), pack-reused 0  
Receiving objects: 100% (3/3), done.  
christian@workstation:~$
```

these two steps are making the repository and cloning the repository in the workstation

```
christian@workstation: ~/CPE232_Bernardo_ACT-10.1
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ sudo nano ansible.cfg
[sudo] password for christian:
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ansible.cfg
[defaults]

inventory = inventory
host_key_checking = false

deprecation_warnings = false

remote_user = christian
private_key_files = ~/.ssh/id_rsa
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

this is the ansible.cfg for the activity

step two: making the inventory file

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ sudo nano inventory
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat inventory
[UbuntuServer]
192.168.56.109

[CentOSServer]
192.168.56.113
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ tree
.
├── ansible.cfg
├── inventory
└── README.md

0 directories, 3 files
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

this is the inventory file

step three: making the directories

```
christian@workstation: ~/CPE232_Bernardo_ACT-10.1
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ mkdir roles
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cd roles
christian@workstation:~/CPE232_Bernardo_ACT-10.1/roles$ mkdir CentOS Ubuntu
christian@workstation:~/CPE232_Bernardo_ACT-10.1/roles$ mkdir ./CentOS/tasks
christian@workstation:~/CPE232_Bernardo_ACT-10.1/roles$ mkdir ./Ubuntu/tasks
christian@workstation:~/CPE232_Bernardo_ACT-10.1/roles$ cd ..
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ tree
.
├── ansible.cfg
├── inventory
├── README.md
└── roles
    ├── CentOS
    │   └── tasks
    └── Ubuntu
        └── tasks

5 directories, 3 files
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

this is the newly made directories including: roles, CentOS, Ubuntu and tasks for each step four: making elst.yml it stands for installation of Elastic Stack

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ sudo nano elst.yml
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat elst.yml
---
- hosts: all
  become: true
  pre_tasks:

    - name: update repository index (CentOS)
      tags: always
      dnf:
        update_cache: yes
        changed_when: false
        when: ansible_distribution == "CentOS"

    - name: install updates (Ubuntu)
      tags: always
      apt:
        update_cache: yes
        changed_when: false
        when: ansible_distribution == "Ubuntu"

- hosts: CentOSServer
  become: true
  roles:
    - CentOS

- hosts: UbuntuServer
  become: true
  roles:
    - Ubuntu
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

step five: the main.yml files for Ubuntu and CentOS

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ sudo nano ./roles/CentOS/tasks/main.yml
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/CentOS/tasks/main.yml
# Elastic Search Setup
```

- name: Temporarily setting the SELINUX of CentOS remote server to permissive
selinux:
 policy: targeted
 state: permissive
 when: ansible_os_family == 'RedHat'
- name: Updating sysctl for max_map_count
sysctl:
 name: vm.max_map_count
 value: "262144"
 sysctl_set: yes
- name: Adding the user 'elasticsearch'
user:
 name: elasticsearch
 comment: elasticsearch user
- name: Creating directory for the downloaded files
file:
 path: /data
 state: directory
 mode: 0777
- name: Extracting elasticsearch
unarchive:
 src: /data/elasticsearch-6.8.15.tar.gz
 dest: /data/
 remote_src: yes
 creates: /data/elasticsearch-6.8.15/config/elasticsearch.yml
- name: Inserting the Elastic Search systemd service unit file
template:
 src: elasticsearch.service.j2
 dest: /etc/systemd/system/elasticsearch.service
 mode: 0644
- name: Inserting the Elastic Search configuration template
template:
 src: elasticsearch.yml.j2
 dest: /data/elasticsearch-6.8.15/config/elasticsearch.yml
 mode: 0660
- file:
 path: /data/elasticsearch-6.8.15
 owner: elasticsearch
 group: elasticsearch
 recurse: yes
- name: Daemon Reload
systemd:
 daemon_reload: yes

```
- name: Starting the Elastic Search service
  service:
    name: elasticsearch
    state: started
    enabled: yes
```

#Kibana Installation and Configuration

```
- name: Creating directory for downloaded files
  file:
    path: /data
    state: directory
    mode: 0777

- name: Installing Kibana tar
  get_url:
    url: https://artifacts.elastic.co/downloads/kibana/kibana-6.8.15-linux-x86_64.tar.gz
    dest: /data/kibana-6.8.15-linux-x86_64.tar.gz
    mode: 0755

- name: Extracting Kibana
  unarchive:
    src: /data/kibana-6.8.15-linux-x86_64.tar.gz
    dest: /data/
    remote_src: yes
    creates: /data/kibana-6.8.15-linux-x86_64/config/kibana.yml
```



```
- name: Inserting the Kibana systemd service unit file
  template:
    src: kibana.service.j2
    dest: /etc/systemd/system/kibana.service
    mode: 0644

- name: Inserting the update of configuration template for Kibana
  template:
    src: kibana.yml.j2
    dest: /data/kibana-6.8.15-linux-x86_64/config/kibana.yml
    mode: 0660

- name: Daemon Reload
  systemd:
    daemon_reload: yes

- name: Starting the Kibana service
  service:
    name: kibana
    state: started
    enabled: yes

# Logstash Setup
- name: Creating directory for downloaded files
  file:
    path: /data
    state: directory
    mode: "u=rwx,g=rwx,o=rwx"
```

```
- name: Installing logstash tar ball
  get_url:
    url: https://artifacts.elastic.co/downloads/logstash/logstash-6.8.15.tar.gz
    dest: /data/logstash-6.8.15.tar.gz
    mode: 0755

- name: Extracting logstash
  unarchive:
    src: /data/logstash-6.8.15.tar.gz
    dest: /data/
    remote_src: yes
    creates: /data/logstash-6.8.15/conf.d/inputs.conf

- name: Inserting the Logstash systemd service unit file
  template:
    src: logstash.service.j2
    dest: /etc/systemd/system/logstash.service
    mode: 0644

- name: Script of logstash for starting/stopping
  template:
    src: start.sh.j2
    dest: /data/logstash-6.8.15/start.sh
    mode: 0754
```

```
- name: Creating /data/logstash-6.8.15/conf.d directory
  file:
    path: /data/logstash-6.8.15/conf.d
    state: directory
    mode: 0777

- name: Updating the configuration default of logstash
  template:
    src: inputs.conf.j2
    dest: /data/logstash-6.8.15/conf.d/inputs.conf
    mode: 0660

- name: Daemon Reload
  systemd:
    daemon_reload: yes

- name: Starting the Logstash service
  service:
    name: logstash
    state: started
    enabled: yes
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

this is the CentOS main.yml

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/Ubuntu/tasks/main.yml
# Elastic Search Setup

- name: Temporarily setting the SELINUX of Ubuntu remote server to permissive
  selinux:
    policy: targeted
    state: permissive
    when: ansible_os_family == 'Ubuntu'

- name: Updating sysctl for max_map_count
  sysctl:
    name: vm.max_map_count
    value: "262144"
    sysctl_set: yes

- name: Adding the user 'elasticsearch'
  user:
    name: elasticsearch
    comment: elasticsearch user

- name: Creating directory for the downloaded files
  file:
    path: /data
    state: directory
    mode: 0777

- name: Downloading elasticsearch tar ball
  get_url:
    url: https://artifacts.elastic.co/downloads/elasticsearch/elasticsearch-6.8.15.tar.gz
    dest: /data/elasticsearch-6.8.15.tar.gz
    mode: 0755
```

```
- name: Extracting elasticsearch
  unarchive:
    src: /data/elasticsearch-6.8.15.tar.gz
    dest: /data/
    remote_src: yes
    creates: /data/elasticsearch-6.8.15/config/elasticsearch.yml

- name: Inserting the Elastic Search systemd service unit file
  template:
    src: elasticsearch.service.j2
    dest: /etc/systemd/system/elasticsearch.service
    mode: 0644

- name: Inserting the Elastic Search configuration template
  template:
    src: elasticsearch.yml.j2
    dest: /data/elasticsearch-6.8.15/config/elasticsearch.yml
    mode: 0660

- file:
    path: /data/elasticsearch-6.8.15
    owner: elasticsearch
    group: elasticsearch
    recurse: yes

- name: Daemon Reload
  systemd:
    daemon_reload: yes
```

```
- name: Starting the Elastic Search service
  service:
    name: elasticsearch
    state: started
    enabled: yes
```

#Kibana Installation and Configuration

```
- name: Creating directory for downloaded files
  file:
    path: /data
    state: directory
    mode: 0777

- name: Installing Kibana tar
  get_url:
    url: https://artifacts.elastic.co/downloads/kibana/kibana-6.8.15-linux-x86_64.tar.gz
    dest: /data/kibana-6.8.15-linux-x86_64.tar.gz
    mode: 0755

- name: Extracting Kibana
  unarchive:
    src: /data/kibana-6.8.15-linux-x86_64.tar.gz
    dest: /data/
    remote_src: yes
    creates: /data/kibana-6.8.15-linux-x86_64/config/kibana.yml
```

```
- name: Inserting the Kibana systemd service unit file
  template:
    src: kibana.service.j2
    dest: /etc/systemd/system/kibana.service
    mode: 0644

- name: Inserting the update of configuration template for Kibana
  template:
    src: kibana.yml.j2
    dest: /data/kibana-6.8.15-linux-x86_64/config/kibana.yml
    mode: 0660

- name: Daemon Reload
  systemd:
    daemon_reload: yes

- name: Starting the Kibana service
  service:
    name: kibana
    state: started
    enabled: yes

# Logstash Setup
- name: Creating directory for downloaded files
  file:
    path: /data
    state: directory
    mode: "u=rwx,g=rwx,o=rwx"
```

```
- name: Installing logstash tar ball
  get_url:
    url: https://artifacts.elastic.co/downloads/logstash/logstash-6.8.15.tar.gz
    dest: /data/logstash-6.8.15.tar.gz
    mode: 0755

- name: Extracting logstash
  unarchive:
    src: /data/logstash-6.8.15.tar.gz
    dest: /data/
    remote_src: yes
    creates: /data/logstash-6.8.15/conf.d/inputs.conf

- name: Inserting the Logstash systemd service unit file
  template:
    src: logstash.service.j2
    dest: /etc/systemd/system/logstash.service
    mode: 0644

- name: Script of logstash for starting/stopping
  template:
    src: start.sh.j2
    dest: /data/logstash-6.8.15/start.sh
    mode: 0754
```

```
- name: Creating /data/logstash-6.8.15/conf.d directory
  file:
    path: /data/logstash-6.8.15/conf.d
    state: directory
    mode: 0777

- name: Updating the configuration default of logstash
  template:
    src: inputs.conf.j2
    dest: /data/logstash-6.8.15/conf.d/inputs.conf
    mode: 0660

- name: Daemon Reload
  systemd:
    daemon_reload: yes

- name: Starting the Logstash service
  service:
    name: logstash
    state: started
    enabled: yes
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

this is the Ubuntu main.yml

step six: Including the required dependencies for Elastic Stack (Elastic Search, Kibana, Logstash)

this is the CentOS dependencies

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ sudo nano ./roles/CentOS/elasticsearch.service.j2
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/CentOS/elasticsearch.service.j2
[Unit]
Description=Elasticsearch service
After=network.target

[Service]
Type=simple
LimitNOFILE=65536
LimitMEMLOCK=infinity
User=elasticsearch
Group=elasticsearch
ExecStart=/data/elasticsearch-6.8.15/bin/elasticsearch

[Install]
WantedBy=multi-user.target
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/CentOS/elasticsearch.yml.j2
http.port: 9200
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ sudo nano ./roles/CentOS/inputs.conf.j2
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/CentOS/inputs.conf.j2
input {
  beats {
    port => 5044
  }
}

filter {
  mutate {
    remove_field => [ '[host]' ]
  }
}

filter {
  mutate {
    convert => { "[system][process][cpu][total][norm][pct]" => "float" }
    convert => { "[system][diskio][iostat][request][avg_size]" => "float" }
    convert => { "[system][process][memory][rss][pct]" => "float" }
    convert => { "[system][process][cpu][total][pct]" => "float" }
    convert => { "[system][diskio][iostat][queue][avg_size]" => "float" }
    convert => { "[system][core][steal][pct]" => "float" }
    convert => { "[system][cpu][steal][pct]" => "float" }
    convert => { "[system][diskio][iostat][service_time]" => "float" }
  }
}

output {
  elasticsearch {
    hosts => "localhost:9200"
    index => 'logstash-daily-%{+YYYY.MM.dd}'
  }
}
```

christian@workstation:~/CPE232_Bernardo_ACT-10.1\$

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ sudo nano ./roles/CentOS/kibana.service.j2
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/CentOS/kibana.service.j2
[Unit]
Description=Kibana service
After=network.target

[Service]
Type=simple
ExecStart=/data/kibana-6.8.15-linux-x86_64/bin/kibana

[Install]
WantedBy=multi-user.target
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ sudo nano ./roles/CentOS/kibana.yml.j2
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/CentOS/kibana.yml.j2
server.port: 5601
server.host: 0.0.0.0
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```



```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ sudo nano ./roles/CentOS/logstash.service.j2
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/CentOS/logstash.service.j2
[Unit]
Description=Logstash service
After=network.target

[Service]
Type=simple
ExecStart=/data/logstash-6.8.15/start.sh

[Install]
WantedBy=multi-user.target
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

this is the Ubuntu dependencies

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/Ubuntu/elasticsearch.service.j2
[Unit]
Description=Elasticsearch service
After=network.target

[Service]
Type=simple
LimitNOFILE=65536
LimitMEMLOCK=infinity
User=elasticsearch
Group=elasticsearch
ExecStart=/data/elasticsearch-6.8.15/bin/elasticsearch

[Install]
WantedBy=multi-user.target
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/Ubuntu/elasticsearch.yml.j2
http.port: 9200
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/Ubuntu/inputs.conf.j2
input {
  beats {
    port => 5044
  }
}

filter {
  mutate {
    remove_field => [ '[host]' ]
  }
}

filter {
  mutate {
    convert => { "[system][process][cpu][total][norm][pct]" => "float" }
    convert => { "[system][diskio][iostat][request][avg_size]" => "float" }
    convert => { "[system][process][memory][rss][pct]" => "float" }
    convert => { "[system][process][cpu][total][pct]" => "float" }
    convert => { "[system][diskio][iostat][queue][avg_size]" => "float" }
    convert => { "[system][core][steal][pct]" => "float" }
    convert => { "[system][cpu][steal][pct]" => "float" }
    convert => { "[system][diskio][iostat][service_time]" => "float" }
  }
}

output {
  elasticsearch {
    hosts => "localhost:9200"
    index => 'logstash-daily-%{+YYYY.MM.dd}'
  }
}
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/Ubuntu/kibana.service.j2
[Unit]
Description=Kibana service
After=network.target

[Service]
Type=simple
ExecStart=/data/kibana-6.8.15-linux-x86_64/bin/kibana

[Install]
WantedBy=multi-user.target
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/Ubuntu/kibana.yml.j2
server.port: 5601
server.host: 0.0.0.0
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/Ubuntu/start.sh.j2
#!/bin/bash

/data/logstash-6.8.15/bin/logstash -f /data/logstash-6.8.15/conf.d/inputs.conf
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ cat ./roles/CentOS/start.sh.j2
#!/bin/bash

/data/logstash-6.8.15/bin/logstash -f /data/logstash-6.8.15/conf.d/inputs.conf
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```

step seven: executing the playbook

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ ansible-playbook --ask-become-pass elst.yml
BECOME password:

PLAY [all] *****

TASK [Gathering Facts] *****
ok: [192.168.56.113]
fatal: [192.168.56.109]: UNREACHABLE! => {"changed": false, "msg": "Failed to connect to the host via ssh: ssh:
connect to host 192.168.56.109 port 22: No route to host", "unreachable": true}

TASK [update repository index (CentOS)] *****
ok: [192.168.56.113]

TASK [install updates (Ubuntu)] *****
skipping: [192.168.56.113]

PLAY [CentOSServer] *****

TASK [Gathering Facts] *****
ok: [192.168.56.113]

TASK [CentOS : Temporarily setting the SELINUX of CentOS remote server to permissive] *****
changed: [192.168.56.113]

TASK [CentOS : Updating sysctl for max_map_count] *****
changed: [192.168.56.113]

TASK [CentOS : Adding the user 'elasticsearch'] *****
changed: [192.168.56.113]

TASK [CentOS : Creating directory for the downloaded files] *****
changed: [192.168.56.113]

TASK [CentOS : Downloading elasticsearch tar ball] *****
changed: [192.168.56.113]
```

```
TASK [CentOS : Extracting elasticsearch] *****
changed: [192.168.56.113]

TASK [CentOS : Inserting the Elastic Search systemd service unit file] *****
changed: [192.168.56.113]

TASK [CentOS : Inserting the Elastic Search configuration template] *****
changed: [192.168.56.113]

TASK [CentOS : file] *****
changed: [192.168.56.113]

TASK [CentOS : Daemon Reload] *****
ok: [192.168.56.113]

TASK [CentOS : Starting the Elastic Search service] *****
changed: [192.168.56.113]

TASK [CentOS : Creating directory for downloaded files] *****
ok: [192.168.56.113]

TASK [CentOS : Installing Kibana tar] *****
changed: [192.168.56.113]

TASK [CentOS : Extracting Kibana] *****
changed: [192.168.56.113]

TASK [CentOS : Inserting the Kibana systemd service unit file] *****
changed: [192.168.56.113]

TASK [CentOS : Inserting the update of configuration template for Kibana] *****
changed: [192.168.56.113]

TASK [CentOS : Daemon Reload] *****
ok: [192.168.56.113]
```

```
TASK [CentOS : Starting the Kibana service] *****
changed: [192.168.56.113]

TASK [CentOS : Creating directory for downloaded files] *****
ok: [192.168.56.113]

TASK [CentOS : Installing logstash tar ball] *****
changed: [192.168.56.113]

TASK [CentOS : Extracting logstash] *****
changed: [192.168.56.113]

TASK [CentOS : Inserting the Logstash systemd service unit file] *****
changed: [192.168.56.113]

TASK [CentOS : Script of logstash for starting/stopping] *****
changed: [192.168.56.113]
```

```

TASK [CentOS : Creating /data/logstash-6.8.15/conf.d directory] *****
changed: [192.168.56.113]

TASK [CentOS : Updating the configuration default of logstash] *****
changed: [192.168.56.113]

TASK [CentOS : Daemon Reload] *****
ok: [192.168.56.113]

TASK [CentOS : Starting the Logstash service] *****
changed: [192.168.56.113]

PLAY [UbuntuServer] *****

PLAY RECAP *****
192.168.56.109      : ok=0    changed=0    unreachable=1    failed=0    skipped=0    rescued=0    ignore
d=0
192.168.56.113    : ok=30   changed=22   unreachable=0    failed=0    skipped=1    rescued=0    ignore
d=0
christian@workstation:~/CPE232_Bernardo_ACT-10.1$

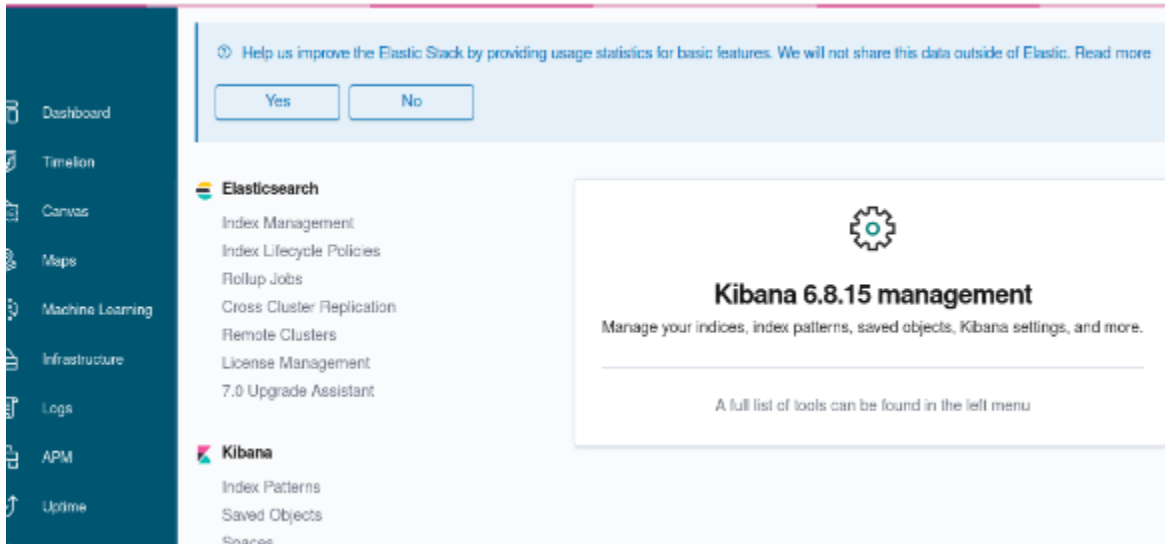
```

step eight: testing the software

this is the CentOS testing

The screenshot shows the Kibana monitoring dashboard in a web browser. The address bar indicates the URL is `192.168.56.113:5601/app/monitoring#/overview?_g=(cluster_uuid:'9LKQQ-`. The browser tabs include 'Centos', 'Wiki', 'Documentation', and 'Forums'. The Kibana sidebar on the left lists various tools like Discover, Visualize, Dashboard, Timeline, Canvas, Maps, Machine Learning, Infrastructure, Logs, APM, Uptime, Dev Tools, Monitoring, and Management. The main content area displays the 'Overview' for the 'elasticsearch' cluster. At the top, there's a message about improving the Elastic Stack by providing usage statistics, with 'Yes' and 'No' buttons. Below this, the 'Elasticsearch' section shows 'Health is green' and 'Basic license'. It contains three summary cards: 'Overview' (Version: 6.8.15, Uptime: an hour), 'Nodes: 1' (Disk Available: 63.43%, JVM Heap: 28.78%), and 'Indices: 4' (Documents: 850, Disk Usage: 981.3 KB). The 'Kibana' section shows 'Health is green' and contains two summary cards: 'Overview' (Requests: 1, Max. Response Time: 869 ms) and 'Instances: 1' (Connections: 0, Memory Usage: 10.16%).

this is the Ubuntu Testing



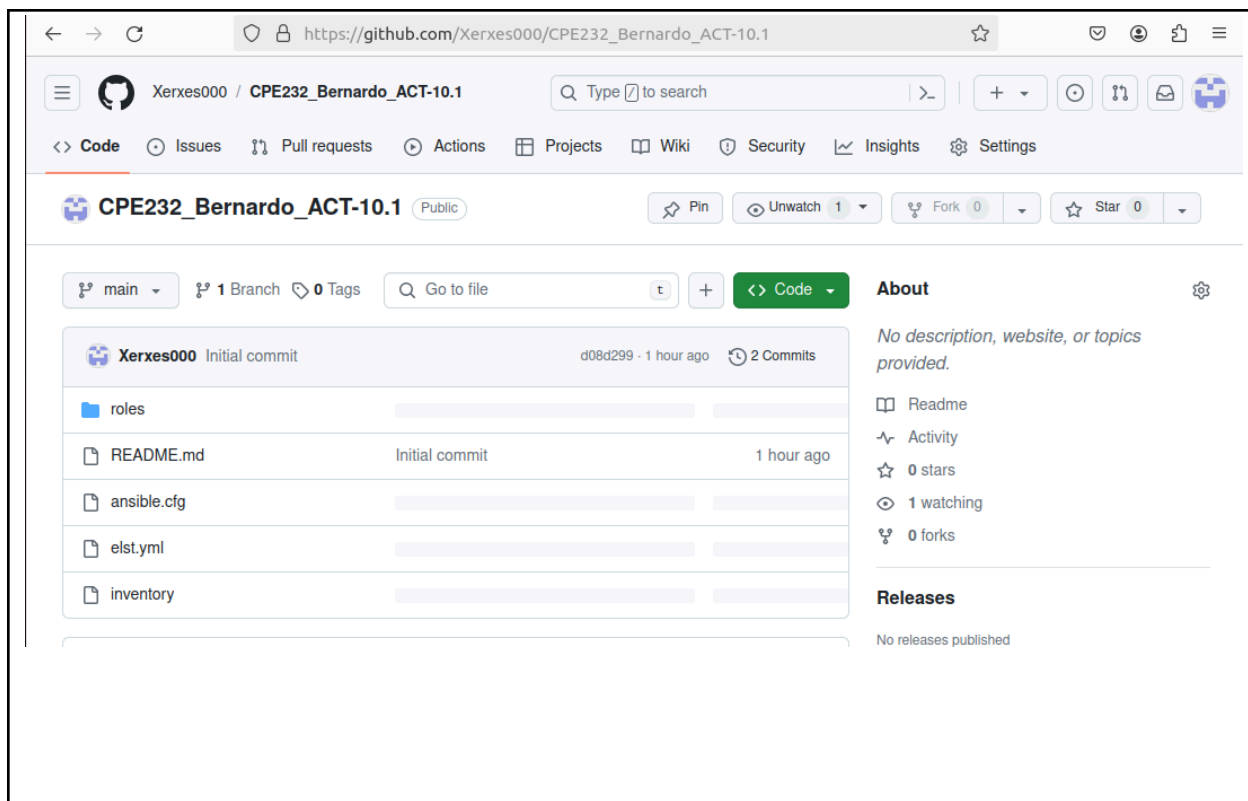
step nine: git pushing the files to the repositories

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ git status
On branch main
Your branch is up to date with 'origin/main'.

Untracked files:
  (use "git add <file>..." to include in what will be committed)
        ansible.cfg
        elst.yml
        inventory
        roles/

nothing added to commit but untracked files present (use "git add" to track)
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ git add *
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ git commit -m "update"
[main 85de48b] update
19 files changed, 509 insertions(+)
create mode 100644 ansible.cfg
create mode 100644 elst.yml
create mode 100644 inventory
create mode 100644 roles/CentOS/elasticsearch.service.j2
create mode 100644 roles/CentOS/elasticsearch.yml.j2
create mode 100644 roles/CentOS/inputs.conf.j2
create mode 100644 roles/CentOS/kibana.service.j2
create mode 100644 roles/CentOS/kibana.yml.j2
create mode 100644 roles/CentOS/logstash.service.j2
create mode 100644 roles/CentOS/start.sh.j2
create mode 100644 roles/CentOS/tasks/main.yml
create mode 100644 roles/Ubuntu/elasticsearch.service.j2
create mode 100644 roles/Ubuntu/elasticsearch.yml.j2
create mode 100644 roles/Ubuntu/inputs.conf.j2
create mode 100644 roles/Ubuntu/kibana.service.j2
create mode 100644 roles/Ubuntu/kibana.yml.j2
create mode 100644 roles/Ubuntu/logstash.service.j2
create mode 100644 roles/Ubuntu/start.sh.j2
create mode 100644 roles/Ubuntu/tasks/main.yml
```

```
christian@workstation:~/CPE232_Bernardo_ACT-10.1$ git push origin main
Enumerating objects: 21, done.
Counting objects: 100% (21/21), done.
Compressing objects: 100% (16/16), done.
Writing objects: 100% (20/20), 3.07 KiB | 629.00 KiB/s, done.
Total 20 (delta 3), reused 0 (delta 0), pack-reused 0
remote: Resolving deltas: 100% (3/3), done.
To github.com:Xerxes000/CPE232_Bernardo_ACT-10.1.git
d08d299..85de48b  main -> main
christian@workstation:~/CPE232_Bernardo_ACT-10.1$
```



- **Github Link:**
https://github.com/Xerxes000/CPE232_Bernardo_ACT-10.1.git

Reflections:

Answer the following:

1. What are the benefits of having log monitoring tool?
 - Workload logs are gathered and centralized logging capabilities are provided by the log monitoring tool. Logs from complicated cloud environments can be combined into a searchable index by using certain log monitoring tools, like the ELK stack. User and network administrators can carry out root cause analysis and enhance security with these features.

Conclusions:

- To summarize, this exercise offered significant understanding into the playbook-based installation and setup of the Elastic Stack, which consists of Kibana, Logstash, and Elastic Search. My practical experience in managing these crucial components came from implementing roles to consolidate playbooks. Creating useful Ansible playbooks that can be used for a variety of tasks was another new skill I picked up. I also gained a deeper understanding of playbook design—more specifically, log monitoring tool design—thanks to this exercise. I now recognize the convenience with which playbooks make the

installation, configuration, and use of these tools possible. And finally, I realized how much benefit network administrators, users, and companies can receive from log monitoring tools. This increases overall system reliability by improving security and facilitating effective root cause analysis.